

Exploring Data with R and Spark

2017 Big Data and Analytics Summer School

Chris Park

2 August 2017

Contents

1	Preface	1
2	Getting Started	1
3	“Programming” in R	2
3.1	Containers and Values	3
3.2	Functions	7
3.3	Quirks	8
4	Exploring tabular data in R	18
4.1	Data Manipulation with <code>dplyr</code>	18
4.2	Data Visualization with <code>ggplot2</code>	38
4.3	“Distributed Processing” with <code>sparklyr</code>	66
4.4	Review	78
4.5	Appendix	79
4.6	R References	79

1 Preface

The aim of today’s workshop is to provide you with hands-on experience in dealing with *reasonably large* data sets using R. By the end of the workshop, you should be able to:

- Read and write simple programs in the R language,
- Manipulate tabular data sets in R using the `dplyr` package,
- Visualize tabular data sets in R using the `ggplot2` package,
- Understand R’s limitations when it comes to large data sets, and
- Manipulate *reasonably large* tabular data sets in R using the `sparklyr` package.

Please drop me an email at chris.park@protonmail.com if you have any comments or suggestions.

2 Getting Started

We will be using [R Notebooks](#) to cover the course material, so that code is interleaved with explanations (which hopefully makes it easier to understand code). These are like [Zeppelin Notebooks](#), but tailor-made for R. Blocks of code within an R Notebook is called a *chunk*, and is visually marked by a grey rectangle. To run the code inside a chunk in RStudio, you can either:

- Click the *Run* button (small green “play” button), or
- Place your cursor inside the chunk and press *Ctrl + Shift + Enter*.

```
## Print greeting message.  
print("Welcome to the 2017 Essex Big Data and Analytics Summer School!")
```

```
## [1] "Welcome to the 2017 Essex Big Data and Analytics Summer School!"
```

It is also possible to insert a new code chunk. To do this, you can either:

- Click the *Insert Chunk* button, or
- Place your cursor inside the chunk and press *Ctrl + Alt + I*.

3 “Programming” in R

A *computer program* is simply a sequence of instructions that describe how we want to perform computations, e.g. add numbers, display messages, and plot graphs. There are many *programming languages* we can choose from to express our computations to the computer. Today, we use the R programming language to express our computations, as it is one of the most widely used languages when it comes to computing with tabular data sets. Using a language to describe our data processing tasks can be very useful:

- More flexibility for handling edge cases e.g. inconsistent formatting,
- Easier to share and reproduce work,
- Automate repetitive tasks instead of copying and pasting multiple times,
- Access to cutting-edge tools developed by the software community,
- etc.

Unlike humans however, computers can’t handle ambiguity, e.g. as high as a kite. They are harsh taskmasters, and don’t tolerate inaccuracies or inconsistencies in instructions. This means that the programs we write must be 100% accurate, and in the *form* that the computer expects. This domain has a big vocabulary so we’re going to have to learn a few words:

Term	Definition
Variable	Containers for values that can change, i.e. are <i>variable</i> .
Expression	Combination of values, variables, and operators.
Assignment	A statement that <i>assigns</i> a <i>value</i> to a <i>variable</i> .
Comment	Notes for humans to read and understand a program.
Function	A <i>named</i> sequence of statements that performs a computation.

For example, in the **expression** below, we say that we are **assigning** the numeric value 1 to the **variable** **a**, rather than **a equals 1**:

```
a = 1
```

Did anything happen? We can check by typing in **a** to see if R returns the value 1.

```
a # Click "Run" or the green triangle or hit "Ctrl + Enter"
```

```
## [1] 1
```

Now, since **a** is a *variable*, we can replace its value by assigning it a new value.

```
a = "Boo!"  
a # 'a' no longer returns the value 1.
```

```
## [1] "Boo!"
```

Note: the sentence **'a' no longer returns the value 1.** in the code above what is called a *comment*. It is just a helpful message for human readers, and ignored by the computer. Any text that follows **#** becomes a *comment* in R. Different programming languages have different symbols that mark the beginning of a comment, e.g. **//** in C and C++.




Figure 1: A **vector** (left) and **matrix** (right)

3.1 Containers and Values

R was written by statisticians, for statisticians. Statisticians often work with *collections* of values rather than individual values, e.g. a survey of annual income from a region rather than a single salary from an individual. For this reason, R was designed to be *collection-oriented*, and the basic data type in R is a **vector**, which is a kind of container that holds collection of values that are all of the *same type*. Broadly speaking, there are four basic *families of values* in R:

- Logical: TRUE, FALSE
- Numeric: 1, 2, 2.0, 3.0, ...
- Character: "Alice", 'Bob', "Cat!"
- Missing: NA (Stands for Not Available)

The four basic types of *containers* for values are:

- Vectors: 1-dimensional collection of items which are of the *same* type.
- Matrices: 2-dimensional collections of items which are of the *same* type.
- Lists: tree-like collection of items where each *item* can be *different*.
- Data Frames: spreadsheet-like collection of items where each *column* can have *different* types of values.

Aside: Recall that R is a *collection-oriented* language designed for statistical computations on *collections* of values, and the basic data type is a **vector**. This means that even a single numeric value like 1 is in fact a **vector** of length 1, rather than a scalar. That is, 1 is identical to `c(1)` in R. This can have some performance implications in complex computations, as R needs to extract the value in and out of a container (i.e. the **vector**) rather than use it directly. This is not the case in most other programming languages.

3.1.1 Vectors

To create a vector, we can combine or concatenate values as follows:

```
## Numeric Vector
c(1, 2)
```

```
## [1] 1 2
```

```
1:2
```

```
## [1] 1 2
```

```
## Character Vector
c("Alice", "Bob", "Cat")

## [1] "Alice" "Bob"  "Cat"
c("A", "B", "C")

## [1] "A" "B" "C"
LETTERS[c(1, 2, 3)]

## [1] "A" "B" "C"
LETTERS[1:3]

## [1] "A" "B" "C"
letters[1:3]

## [1] "a" "b" "c"
## Logical Vector
c(TRUE, FALSE)

## [1] TRUE FALSE
c(1, 2)[c(TRUE, FALSE)]

## [1] 1
## In fact, even the single numeric value 1 is a vector.
1

## [1] 1
is.vector(1)

## [1] TRUE
## Exercise 1.1: Explain what is going on in the code below:
c(1, "A")

## [1] "1" "A"
c(1, 2, 3, 4)[c(TRUE, FALSE)]

## [1] 1 3
c(1, 2, 3, 4)[c(TRUE, FALSE, TRUE, FALSE)]

## [1] 1 3
```

3.1.2 Matrices

```
## Numeric Matrix
i = 1:6
matrix(i, nrow = 3)

##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
matrix(i, nrow = 3, ncol = 2)
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
matrix(i, ncol = 2)
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
cbind(1:3, 4:6) # column bind
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
rbind(c(1, 4), c(2, 5), c(3, 6)) # row bind
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
## Character Matrix
```

```
matrix(LETTERS[i], nrow = 3)
```

```
##      [,1] [,2]
## [1,] "A"  "D"
## [2,] "B"  "E"
## [3,] "C"  "F"
```

```
matrix(i, nc = 2, byrow = TRUE)
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
## [3,]    5    6
```

3.1.3 Lists

```
x1 = list(1, "Alice")
x2 = list(TRUE, 2, "Bob")
x3 = list(x1, x2)
x3
```

```
## [[1]]
## [[1]][[1]]
## [1] 1
##
## [[1]][[2]]
## [1] "Alice"
##
```




Figure 2: A list (left) and `data.frame` (right)

```
##
## [[2]]
## [[2]][[1]]
## [1] TRUE
##
## [[2]][[2]]
## [1] 2
##
## [[2]][[3]]
## [1] "Bob"
```

```
x1
```

```
## [[1]]
## [1] 1
##
## [[2]]
## [1] "Alice"
```

3.1.4 Data Frames

```
x1 = 1:3
x2 = LETTERS[x1]
x3 = data.frame(id = x1, cat = x2)
x3
```

```
##   id cat
## 1  1  A
## 2  2  B
## 3  3  C
```

```
data.frame(x3, x3)
```

```
##   id cat id.1 cat.1
```

```
## 1 1 A 1 A
## 2 2 B 2 B
## 3 3 C 3 C

data.frame(x3, x3, x3)

##   id cat id.1 cat.1 id.2 cat.2
## 1 1 A 1 A 1 A
## 2 2 B 2 B 2 B
## 3 3 C 3 C 3 C

data.frame(list(a = x1, b = x2))

##   a b
## 1 1 A
## 2 2 B
## 3 3 C
```

3.1.5 Checking Types

```
is.numeric(1:10)

## [1] TRUE

is.numeric(c("A", "B"))

## [1] FALSE

is.data.frame(matrix(1:10, nrow = 2))

## [1] FALSE

is.matrix(data.frame(1:10, 2:11))

## [1] FALSE

is.list(matrix(1:10, ncol = 2))

## [1] FALSE

## Exercise 1.2: is a data.frame a list?

## Exercise 1.3: is a list a data.frame?
```

3.2 Functions

Functions allow us to avoid writing the same chunk of code over and over again. We can use them to package up a group of statements into a single bundle, and give it a name of our choice. It is good practice to give it a name that is related to what it was written for. For example, the `print()` function prints a message on screen. Suppose we wanted to create a kind of greeting system to `print()` the following greeting message to students.

```
## Print a greeting message for summer school students.
print("Hello, Peter! Welcome to the 2017 Essex BDAS!")

## [1] "Hello, Peter! Welcome to the 2017 Essex BDAS!"

print("Hello, Chris! Welcome to the 2017 Essex BDAS!")
```

```
## [1] "Hello, Chris! Welcome to the 2017 Essex BDAS!"
print("Hello, Sarah! Welcome to the 2017 Essex BDAS!")
```

```
## [1] "Hello, Sarah! Welcome to the 2017 Essex BDAS!"
```

Suppose we wanted to `print()` the message to 200 students. We will then have to repeat ourselves an awful lot, as the only portion of the code that needs changing is the name of each student. It would be very useful to have a way of simply slotting in a student's name into the message. This is exactly what a **function** allows us to do:

```
## We define a function called "greet_student" that prints a friendly
## welcome message to 2017 Essex BDAS students. It takes a single argument
## called "name", which should be a character vector.
greet_student = function(name) {
  print(paste0("Hello, ",
               name,
               "! ",
               "Welcome to the 2017 Essex BDAS!"))
}

## "sapply()" applies a function to each element in a vector.
student_names = c("Simon", "Jack", "Helen", "Alex")
sapply(student_names, greet_student)
```

```
## [1] "Hello, Simon! Welcome to the 2017 Essex BDAS!"
## [1] "Hello, Jack! Welcome to the 2017 Essex BDAS!"
## [1] "Hello, Helen! Welcome to the 2017 Essex BDAS!"
## [1] "Hello, Alex! Welcome to the 2017 Essex BDAS!"

##                               Simon
## "Hello, Simon! Welcome to the 2017 Essex BDAS!"
##                               Jack
## "Hello, Jack! Welcome to the 2017 Essex BDAS!"
##                               Helen
## "Hello, Helen! Welcome to the 2017 Essex BDAS!"
##                               Alex
## "Hello, Alex! Welcome to the 2017 Essex BDAS!"
```

3.3 Quirks

3.3.1 Flexibility vs Speed

R is a very *flexible* language, and the value of almost anything can be modified on the fly. This means that R has to look up the meaning of a value every time it performs a computation. Imagine you're reading a book, and the words on the page are constantly changing into a different language; take the word "zero" for example: zero (English, Italian, etc.), teg (Mongolian), cero (Spanish), nul (Danish), etc. This is going to slow down your reading speed as you would have to look up the meaning of every word as you read. This is kind of what R is doing in the background.

In general, there is a speed-flexibility trade-off when it comes to programming languages: inflexible languages tend to be faster, while flexible languages tend to be slower. For example, in an inflexible language like C or C++, you have to specify the types of every value you use in a program. For example, suppose we wanted to compute the column sums of a matrix.


```
## This chunk has been commented out deliberately to save time.
## Feel free to uncomment this by highlighting this chunk and
## pressing "Ctrl + Shift + C"
```

```
# library(Rcpp)
#
# ## Observe that we declare the variable and function types.
# cppFunction('NumericVector colSums_C(NumericMatrix x) {
#   int nrow = x.nrow(), ncol = x.ncol();
#   NumericVector out(ncol);
#   for (int j = 0; j < ncol; j++) {
#     double total = 0;
#     for (int i = 0; i < nrow; i++) {
#       total += x(i, j);
#     }
#     out[j] = total;
#   }
#   return out;
# }')
#
# colSums_R = function(x) {
#   n = ncol(x)
#   out = numeric(n)
#   for (j in 1:n) {
#     out[j] = sum(x[, j])
#   }
#   out
# }
#
# ## Simulate a 10 million by 2 matrix of random normals.
# A = cbind(rnorm(1e8), rnorm(1e8))
#
# ## Compare times.
# system.time(colSums_R(A))
# # user system elapsed
# # 1.90    0.14    2.05
#
# system.time(colSums_C(A))
# # user system elapsed
# # 0.32    0.00    0.33
```

As we can see above, we don't need to declare types at all with R – we can simply define things on the fly; but this flexibility comes at the cost of speed because R is constantly looking things up in its dictionary, so it's doing a lot of work in the background. This kind of behaviour can also lead to obscure inconsistencies in your code and make it difficult to debug programs (this becomes particularly problematic with larger datasets).

```
c(1, 2, 3, 4)

c = sum
c(1, 2, 3, 4)

rm(c)
```

We can even do crazy things like:



Figure 3: A book within a book to depict R's scoping rules

```
2 * 2
2 / 2

`*` = `/`
rm = `+`

2 * 2
2 / 2

rm(`*`)

## We can also do crazy things like rm = sum!
```

Things can get even more interesting/confusing when you take into account *how* R looks up the meaning of things: R looks up the meaning of words from the inner-most dictionary first. This is what happens in the code below - defining a new *function* results in a new *inner dictionary* from which to look up the meaning of objects from.

```
## Exercise 1.4: What is happening in the code below?
f = function(n) {
  `*` = `/`
  n * n
}

f(2)

## [1] 1
```

```

2 * 2

## [1] 4

## Bonus
f = function(n) {
  g = function(n) {
    `*` = `/`
    n * n
  }
  g(n)
}
f(2)

## [1] 1

## Flip a coin and define `*` to be `/` if it turns up heads.
f = function() {
  if (runif(1) > .5) {
    `*` = `/`
  }
  1 * 1
}

## Run the above 10000 times and see what happens.
iter = 1e5
res = numeric(iter)

## Approach 1
system.time({
  for (i in 1:iter) {
    res[i] = f()
  }
})

##      user  system elapsed
##    0.41    0.00    0.41

## user  system elapsed
## 0.28    0.00    0.29

res = f()
table(res)

## res
## 1
## 1

## Approach 2
system.time({
  for (i in 1:iter) {
    res = c(res, f())
  }
})

##      user  system elapsed
##   24.09    0.34   24.51

```

```
## user system elapsed
## 14.02    0.02    14.04
```

3.3.2 Copying

R also makes a lot of copies of objects. This seems to be historical baggage from the early days of statistical software, when software was developed on machines with very limited RAM (e.g. 32KB – compare that to 32GB of RAM on my machine, which is 32 million KB!). This meant that data couldn't fit into memory so small chunks of the data had to be *copied* into memory at each stage of the analysis. Unfortunately, this can lead to lots of unnecessary copying, which can be particularly problematic with big data sets, especially if they are stored in inefficient data structures like data frames. Data frames can cause problems as even a simple update leads to multiple copies of the data frame being made.

```
## We create a data frame consisting of 2 million random integers.
## Imagine that these are counts the number of steps people have
## taken across two days.
day1 = sample(9000:20000, 1e6, replace = TRUE)
day2 = sample(9000:20000, 1e6, replace = TRUE)
steps = data.frame(day1, day2)
head(steps)
```

```
##   day1 day2
## 1 15978 13177
## 2 16463 16627
## 3 16222 15785
## 4 13556 17930
## 5 12099 15197
## 6 12588 14530
```

```
## Suppose that there was a fault in the collection device and we need
## to increment the first 1000 counts for both days.
m = 1000
```

```
## Data frame way.
# system.time({
#   for (i in 1:m) {
#     steps[i, ] = steps[i, ] + 1
#   }
# })
# user system elapsed
# 10.35    3.78    14.22
```

```
## Vector way. This offloads the update process down to C code, so
## the copying is under control.
```

```
# system.time({
#   for (i in 1:m) {
#     day1[i] = day1[i] + 1
#     day2[i] = day2[i] + 1
#   }
#   steps2 = data.frame(day1, day2)
# })
# user system elapsed
# 0.02    0.00    0.01
#
```

```
# 14.22 / 0.01 ~ 1400

# identical(steps, steps2)

## Just like that, we achieved almost a 1500-fold increase in speed. While
## this may seem like a trivial exercise, its implications are quite
## profound: having an understanding of how R works under the hood can save
## you a lot of heartache (and money) down the line.
```

3.3.3 “Big” R

R has a physical limit to the size of data sets it can handle. This is because it requires that the data set of interest fits into RAM all at once (note: A 32-bit OS can’t address more than 4GB of memory, and only a fraction of this is made available to R. With 64-bit machines, it all depends on how much money you’re willing to spend on buying additional RAM). We can view our memory limit, in MBs, as follows:

```
memory.limit()

## [1] 32622

Let’s try and hit our memory limit by simulating a large matrix.
```

```
## Simulate a large matrix.
zero_mat = matrix(0, 1e8, ncol = 10)
dim(zero_mat)

## [1] 100000000      10

print(object.size(zero_mat), unit = "Gb")

## 7.5 Gb
```

This time, we simulate a large matrix of *integer* values.

```
## Simulate a large integer matrix.
int_mat = matrix(as.integer(0), 1e+8, ncol = 10)
dim(int_mat)

## [1] 100000000      10

print(object.size(int_mat), unit = "Gb")

## 3.7 Gb
```

Why is the *integer* matrix so much smaller? This is because *numeric* values take up 8 bytes by default, whereas *integer* values only take up 4. This means that if we know we only need to handle *integer* values, specifying this at the beginning when we create a data structure to hold these, can save a lot of memory.

Observe that we’ve ended up repeating ourselves a lot. If we wanted to experiment with different types of simulated matrices, it would be convenient to have a function that does this for us. Writing functions can be a bit tricky in the beginning, but it is important to get used to this way of thinking. Writing *pseudocode* can help you get started.

```
sim_mat =
  function(type, val, nrow, ncol, unit) {
    <<coerce values to specified type>>
    <<simulate large matrix with specified dims>>
    <<print the memory footprint in specified unit>>
  }
```




Figure 4: Parallel Processing

It is the conceptual deconstruction reflected in the pseudocode above that is most important. Translating this into a programming language (say, R) is largely a trivial task.

```
## Write the "sim_mat" function. Notice that we provide default values for
## all the arguments, for convenience.
sim_mat =
  function(type = "numeric", val = 0, nrow = 1e+8, ncol = 10, unit = "Gb") {
    type_val = as(val, type)
    mat = matrix(type_val, nrow, ncol)
    print(object.size(mat), unit = unit)
  }

## We call the sim_mat function with different types of values.
# sim_mat("numeric")
# sim_mat("integer")
```

Now, let's really try and reach out memory limit.

```
# sim_mat(ncol = 1000)
```

3.3.4 Parallel Processing

Now, recall that we talked a bit about *distributed* processing and Map Reduce earlier (in our slides). The basic idea is to break down a big task into small chunks (a bit like programming itself), perform operations on each chunk in parallel (e.g. `Map()`), then combine the result (e.g. `Reduce()`). We can simulate this behaviour within R as follows (note: the example below is over-simplified and strictly for demonstration purposes only. In practice, it would make no sense to do this kind of splitting and combining with such a small vector, and for an operation as simple as adding up a bunch of numeric values).

```
## Split the numeric vector `x` into 4 chunks.
x = 1:100
x_split = split(x, 1:4)
x_split

## $`1`
## [1] 1 5 9 13 17 21 25 29 33 37 41 45 49 53 57 61 65 69 73 77 81 85 89
```

```
## [24] 93 97
##
## $`2`
## [1] 2 6 10 14 18 22 26 30 34 38 42 46 50 54 58 62 66 70 74 78 82 86 90
## [24] 94 98
##
## $`3`
## [1] 3 7 11 15 19 23 27 31 35 39 43 47 51 55 59 63 67 71 75 79 83 87 91
## [24] 95 99
##
## $`4`
## [1] 4 8 12 16 20 24 28 32 36 40 44 48 52 56 60 64 68
## [18] 72 76 80 84 88 92 96 100
```

```
## Compute the sum of each chunk, by "mapping" the function "sum" to each
## chunk in "x_split".
comp = Map(sum, x_split)
comp
```

```
## $`1`
## [1] 1225
##
## $`2`
## [1] 1250
##
## $`3`
## [1] 1275
##
## $`4`
## [1] 1300
```

```
## Combine results, i.e. "Reduce".
Reduce(sum, comp)
```

```
## [1] 5050
```

```
## Identical to:
sum(x)
```

```
## [1] 5050
```

Note that we have used multiple variables to store *intermediate results* (e.g. `x_split`, `comp`). It can sometimes be useful to be able to *chain together* expressions, using the *pipe operator* `%>%`. This operator *pipes* the result of the current expression into another expression. This can make code easier to follow and create less variable clutter. We will be making heavy use of this operator in our workshop today.

```
## Version 1: Intermediate variables
x_split = split(x, 1:4)
comp = Map(sum, x_split)
Reduce(sum, comp)
```

```
## [1] 5050
```

```
## Version 2: Nested calls
Reduce(sum, Map(sum, split(x, 1:4)))
```

```
## [1] 5050
```

```
## Version 3: Pipe operator; arguably more compact and easier to follow.
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
x %>%
  split(., 1:4) %>% # "." is a placeholder for "x".
  Map(sum, .) %>%
  Reduce(sum, .)
```

```
## [1] 5050

## Another benefit is that other mainstream programming languages used in
## data science read similarly. For example, you just have to replace
## `%>%` with say a ``.` operator in Python.
```

We can also use the pipe operator within functions. Suppose we wanted to generalize this Map Reduce example, so that we can split the data into more chunks.

```
sum_MR = function(x = 1:100, n = 4) {
  x %>%
    split(., 1:n) %>% {
      ## Curly braces can be used to group 1+ statements.
      ## Here we print the split dataset for confirmation.
      print(.)
      Map(sum, .)
    } %>%
    Reduce(sum, .)
}

sum_MR()
```

```
## $`1`
## [1] 1 5 9 13 17 21 25 29 33 37 41 45 49 53 57 61 65 69 73 77 81 85 89
## [24] 93 97
##
## $`2`
## [1] 2 6 10 14 18 22 26 30 34 38 42 46 50 54 58 62 66 70 74 78 82 86 90
## [24] 94 98
##
## $`3`
## [1] 3 7 11 15 19 23 27 31 35 39 43 47 51 55 59 63 67 71 75 79 83 87 91
## [24] 95 99
##
## $`4`
## [1] 4 8 12 16 20 24 28 32 36 40 44 48 52 56 60 64 68
## [18] 72 76 80 84 88 92 96 100
## [1] 5050
```



```
sum_MR(n = 5)
```

```
## $`1`  
## [1] 1 6 11 16 21 26 31 36 41 46 51 56 61 66 71 76 81 86 91 96  
##  
## $`2`  
## [1] 2 7 12 17 22 27 32 37 42 47 52 57 62 67 72 77 82 87 92 97  
##  
## $`3`  
## [1] 3 8 13 18 23 28 33 38 43 48 53 58 63 68 73 78 83 88 93 98  
##  
## $`4`  
## [1] 4 9 14 19 24 29 34 39 44 49 54 59 64 69 74 79 84 89 94 99  
##  
## $`5`  
## [1] 5 10 15 20 25 30 35 40 45 50 55 60 65 70 75 80 85  
## [18] 90 95 100  
  
## [1] 5050
```

The above was *simulating* parallel processing to help visualize the process. R also has parallel computing capabilities.

```
# library(parallel)  
#  
# ## How many CPU cores are available?  
# nc = detectCores() - 1  
#  
# ## Create a local cluster of processing nodes.  
# cl = makeCluster(nc)  
#  
# ## Split data into partitions (~ same number in each)  
# pt = clusterSplit(cl, 1:1e8)  
#  
# ## Parallel supply.  
# system.time(parSapply(cl, pt, function(x) x^2))  
# # user system elapsed  
# # 1.47 1.17 3.23  
#  
# ## Serial supply  
# system.time(sapply(pt, function(x) x^2))  
# # user system elapsed  
# # 0.30 0.03 0.33  
#  
# ## Return resources to the OS.  
# stopCluster(cl)
```

Note that running in parallel doesn't always lead to faster execution times. Sometimes, the CPU is multitasking on other processes or has to wait for system resources like network connections to become available. More importantly:

- Communication: data needs to be transferred back and forth between processes.
- Collision: processes can get in each other's way when we try and access the same data at the same time.
- Load balancing: work needs to be divided fairly across the processes. Otherwise, some processes may remain idle and unproductive when there is still lots of work to be done.

All of this adds overhead. Suppose you were tasked with mowing a large lawn. You could get a few of your friends to help out, but it wouldn't make much sense if the lawn was small enough for you to mow by yourself. In this case, it might be better to invest in a better lawn mower. Alternatively, having a few friends help for a reasonably large lawn might make sense, but having 10 won't - you might keep running into each other, it will be hard to communicate who is doing which section of the lawn and keep track of each other's progress, etc.

4 Exploring tabular data in R

4.1 Data Manipulation with dplyr

Disclaimer: figures and exercises in this section are adapted from [Modern Data Science with R](#) by Baumer et al. and [R for Data Science](#) by Wickham & Grolemund.

Now that we're somewhat familiar with the R environment, we will move on to handling some actual data. In this section, we use real data sets from the `nycflights13` package, which contains information about flight arrival and departures for all 336,776 commercial flights in the following airports in New York City in 2013.

Table	Description
<code>airlines</code>	Airline carrier codes and names.
<code>airports</code>	Airport names and locations.
<code>flights</code>	Airline on-time data for all flights departing NYC in 2013.
<code>planes</code>	Construction information about each plane.
<code>weather</code>	hourly weather data for each airport.

The data comes from the US [Bureau of Transportation Statistics](#); detailed documentation can be found in `?nycflights13`.

We can use this dataset to answer questions such as:

- When is the best time to travel?
- Does bad weather cause flight delays?
- Do single delays cause multiple delays?
- Do older planes suffer from more delays?

We will be using the `dplyr` package to explore the data. This package provides 5 *verbs* that can be used for simple data manipulation tasks. These provide a rich means of slicing-and-dicing a single table of data when used in conjunction with each other.

Verb	Description
<code>select</code>	<code>select</code> a subset of columns from a data frame.
<code>filter</code>	<code>filter</code> (i.e. select a subset) rows of a data frame.
<code>arrange</code>	<code>arrange</code> (i.e. reorder) rows in ascending or descending order.
<code>mutate</code>	<code>mutate</code> (i.e. transform) existing columns to derive new columns.
<code>summarize</code>	<code>summarize</code> (i.e. collapse) the rows down to a single row.

The verbs can also be used in conjunction with `group_by` to partition a table into specified *groups* of rows and applying a summary operation by group, e.g. the range of values contained within each group.

All `dplyr` verbs work in a similar way:

- The first *argument* is a data frame,
- The subsequent arguments describe what to do with the data frame.




Figure 5: `The dplyr::select()` function

- The result is a new data frame.

```
## Load the required R packages.
```

```
library(nycflights13)
```

```
library(dplyr)
```

```
## Let's start by looking at the "flights" table.
```

```
head(flights)
```

```
## # A tibble: 6 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     517           515           2     830
```

```
## 2  2013     1     1     533           529           4     850
```

```
## 3  2013     1     1     542           540           2     923
```

```
## 4  2013     1     1     544           545          -1    1004
```

```
## 5  2013     1     1     554           600          -6     812
```

```
## 6  2013     1     1     554           558          -4     740
```

```
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
```

```
## #   time_hour <dtm>
```

```
## How big is it?
```

```
dim(flights) # 336,776 rows by 19 columns, or 6,398,744 observations.
```

```
## [1] 336776      19
```

4.1.1 `select()`

The `flights` table has 19 columns. If we know in advance which columns we are interested in, it would make sense to take a subset of the table by **selecting** a few columns, e.g. departure and arrival times/delays.

```
## Take the table "flights", select the desired columns, and print the
```

```
## dimensions of the table.
```

```
flights %>%
```

```
  select(dep_delay, arr_delay, dep_time, arr_time) %>%
```

```
dim()

## [1] 336776      4

## Exercise 2.1: Create a variable called "times_df" containing the table
##               above, and print the first few rows using "head()"
times_df = flights %>%
  select(dep_delay, arr_delay, dep_time, arr_time)

head(times_df)

## # A tibble: 6 x 4
##   dep_delay arr_delay dep_time arr_time
##   <dbl>     <dbl>   <int>   <int>
## 1         2         11     517     830
## 2         4         20     533     850
## 3         2         33     542     923
## 4        -1        -18     544    1004
## 5        -6        -25     554     812
## 6        -4         12     554     740
```

Sometimes it can be very useful to be able to extract columns by specifying a textual pattern:

```
flights %>%
  select(contains("delay"))

## # A tibble: 336,776 x 2
##   dep_delay arr_delay
##   <dbl>     <dbl>
## 1         2         11
## 2         4         20
## 3         2         33
## 4        -1        -18
## 5        -6        -25
## 6        -4         12
## 7        -5         19
## 8        -3        -14
## 9        -3         -8
## 10       -2          8
## # ... with 336,766 more rows

flights %>%
  select(starts_with("dep"))

## # A tibble: 336,776 x 2
##   dep_time dep_delay
##   <int>     <dbl>
## 1     517         2
## 2     533         4
## 3     542         2
## 4     544        -1
## 5     554        -6
## 6     554        -4
## 7     555        -5
## 8     557        -3
## 9     557        -3
## 10    558        -2
```

```
## # ... with 336,766 more rows
```

```
flights %>%  
  select(ends_with("delay"))
```

```
## # A tibble: 336,776 x 2
```

```
##   dep_delay arr_delay
```

```
##   <dbl>      <dbl>
```

```
## 1         2         11
```

```
## 2         4         20
```

```
## 3         2         33
```

```
## 4        -1        -18
```

```
## 5        -6        -25
```

```
## 6        -4         12
```

```
## 7        -5         19
```

```
## 8        -3        -14
```

```
## 9        -3         -8
```

```
## 10       -2          8
```

```
## # ... with 336,766 more rows
```

```
## Exercise 2.2: Write some R code to return a data frame which only contains
```

```
##               columns that begin with "arr" and end with "delay".
```

```
##               (HINT: try "select()"ing twice.)
```

```
flights %>%  
  select(starts_with("arr")) %>%  
  select(ends_with("delay"))
```

```
## # A tibble: 336,776 x 1
```

```
##   arr_delay
```

```
##   <dbl>
```

```
## 1         11
```

```
## 2         20
```

```
## 3         33
```

```
## 4        -18
```

```
## 5        -25
```

```
## 6         12
```

```
## 7         19
```

```
## 8        -14
```

```
## 9         -8
```

```
## 10          8
```

```
## # ... with 336,766 more rows
```

```
## Equivalent to:
```

```
grep("^arr|delay$", names(flights), value = TRUE)
```

```
## [1] "dep_delay" "arr_time"  "arr_delay"
```

4.1.2 filter()

We can also filter rows by a specified condition. For example, we could extract information about all flights that flow on Christmas day. Note the use of the == operator for checking equality. This is because = is used for variable assignment and function argument specification.

```
## Christmas flights
```

```
flights %>%  
  filter(month == 12, day == 25) # Simply list conditions, separated by ","
```




Figure 6: `Thedplyr::filter()` function

```
## # A tibble: 719 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013    12    25     456           500        -4     649
## 2  2013    12    25     524           515         9     805
## 3  2013    12    25     542           540         2     832
## 4  2013    12    25     546           550        -4    1022
## 5  2013    12    25     556           600        -4     730
## 6  2013    12    25     557           600        -3     743
## 7  2013    12    25     557           600        -3     818
## 8  2013    12    25     559           600        -1     855
## 9  2013    12    25     559           600        -1     849
## 10 2013    12    25     600           600         0     850
## # ... with 709 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

## Flights with a departure delay of more than 2 hours
flights %>%
  filter(dep_delay >= 120)

## # A tibble: 9,888 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     848           1835      853    1001
## 2  2013     1     1     957           733      144    1056
## 3  2013     1     1    1114           900      134    1447
## 4  2013     1     1    1540          1338      122    2020
## 5  2013     1     1    1815          1325      290    2120
## 6  2013     1     1    1842          1422      260    1958
## 7  2013     1     1    1856          1645      131    2212
## 8  2013     1     1    1934          1725      129    2126
## 9  2013     1     1    1938          1703      155    2109
## 10 2013     1     1    1942          1705      157    2124
## # ... with 9,878 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Flights operated by JetBlue (B6) or ExpressJet (EV)
```

```
flights %>%
```

```
  filter(carrier %in% c("B6", "EV"))
```

```
## # A tibble: 108,808 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     544           545          -1    1004
```

```
## 2  2013     1     1     555           600          -5     913
```

```
## 3  2013     1     1     557           600          -3     709
```

```
## 4  2013     1     1     557           600          -3     838
```

```
## 5  2013     1     1     558           600          -2     849
```

```
## 6  2013     1     1     558           600          -2     853
```

```
## 7  2013     1     1     559           559           0     702
```

```
## 8  2013     1     1     600           600           0     851
```

```
## 9  2013     1     1     601           600           1     844
```

```
## 10 2013     1     1     613           610           3     925
```

```
## # ... with 108,798 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
## Winter flights
```

```
flights %>%
```

```
  filter(between(month, 1, 3))
```

```
## # A tibble: 80,789 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     517           515           2     830
```

```
## 2  2013     1     1     533           529           4     850
```

```
## 3  2013     1     1     542           540           2     923
```

```
## 4  2013     1     1     544           545          -1    1004
```

```
## 5  2013     1     1     554           600          -6     812
```

```
## 6  2013     1     1     554           558          -4     740
```

```
## 7  2013     1     1     555           600          -5     913
```

```
## 8  2013     1     1     557           600          -3     709
```

```
## 9  2013     1     1     557           600          -3     838
```

```
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 80,779 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
## Cancelled flights
```

```
flights %>%
```

```
  filter(is.na(dep_time), is.na(arr_time))
```

```
## # A tibble: 8,255 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     NA           1630          NA     NA
```

```
## 2  2013     1     1     NA           1935          NA     NA
```

```
## 3  2013     1     1     NA           1500          NA     NA
```

```
## 4  2013     1     1     NA           600          NA     NA
```

```
## 5 2013 1 2 NA 1540 NA NA
## 6 2013 1 2 NA 1620 NA NA
## 7 2013 1 2 NA 1355 NA NA
## 8 2013 1 2 NA 1420 NA NA
## 9 2013 1 2 NA 1321 NA NA
## 10 2013 1 2 NA 1545 NA NA
## # ... with 8,245 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Exercise 2.3:
```

```
##
```

```
## Which of the flights:
```

```
##
```

```
## 1. Arrived more than three hours late but didn't leave late.
```

```
flights %>%
```

```
  filter(arr_delay > 180, dep_delay <= 0, !is.na(dep_delay))
```

```
## # A tibble: 1 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
```

```
## 1  2013   11     1     658           700        -2    1329
```

```
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
```

```
## #   time_hour <dtm>
```

```
## 2. Flew in January, March, or May.
```

```
flights %>%
```

```
  filter(month %in% c(1, 3, 5))
```

```
## # A tibble: 84,634 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
```

```
## 1  2013     1     1     517           515         2     830
```

```
## 2  2013     1     1     533           529         4     850
```

```
## 3  2013     1     1     542           540         2     923
```

```
## 4  2013     1     1     544           545        -1    1004
```

```
## 5  2013     1     1     554           600        -6     812
```

```
## 6  2013     1     1     554           558        -4     740
```

```
## 7  2013     1     1     555           600        -5     913
```

```
## 8  2013     1     1     557           600        -3     709
```

```
## 9  2013     1     1     557           600        -3     838
```

```
## 10 2013     1     1     558           600        -2     753
```

```
## # ... with 84,624 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
## 3. Were operated by Hawaiian Airlines, Endeavour Air, or
```

```
##   Frontier Airlines and ended up being cancelled.
```

```
##   Look at the 'airlines' table for carrier codes.
```

```
## "HA", "9E", and "F9"
```

```
a1 = flights %>%
```





Figure 7: The `dplyr::arrange()` function

```
## select(carrier) %>%
filter(carrier == c("HA", "9E", "F9"))

## Warning in carrier == c("HA", "9E", "F9"): longer object length is not a
## multiple of shorter object length

a2 = flights %>%
  ## select(carrier) %>%
  filter(carrier %in% c("HA", "9E", "F9"))
```

4.1.3 arrange()

The next verb is `arrange`: we can reorder the rows in ascending or descending order of specified columns. This would be very useful for finding out things like which of the flights had the longest or shortest delay.

```
## Flights with the longest departure delays.
flights %>%
  select(dep_delay) %>%
  arrange(dep_delay)

## # A tibble: 336,776 x 1
##   dep_delay
##   <dbl>
## 1      -43
## 2      -33
## 3      -32
## 4      -30
## 5      -27
## 6      -26
## 7      -25
## 8      -25
## 9      -24
## 10     -24
## # ... with 336,766 more rows

flights %>%
  select(dep_delay) %>%
  arrange(desc(dep_delay))

## # A tibble: 336,776 x 1
```

```
##      dep_delay
##      <dbl>
## 1      1301
## 2      1137
## 3      1126
## 4      1014
## 5      1005
## 6       960
## 7       911
## 8       899
## 9       898
## 10      896
## # ... with 336,766 more rows
```

```
## Fastest flights
flights %>%
  select(air_time) %>%
  arrange(air_time)
```

```
## # A tibble: 336,776 x 1
##   air_time
##   <dbl>
## 1      20
## 2      20
## 3      21
## 4      21
## 5      21
## 6      21
## 7      21
## 8      21
## 9      21
## 10     21
## # ... with 336,766 more rows
```

```
## We can slice() a data frame to extract rows by index. For example,
## we can extract the 15th - 20th fastest flights:
flights %>%
  select(air_time) %>%
  arrange(air_time) %>%
  slice(15:20)
```

```
## # A tibble: 6 x 1
##   air_time
##   <dbl>
## 1      21
## 2      21
## 3      22
## 4      22
## 5      22
## 6      22
```

```
## Notice that there are a few duplicated values for 'air_time'.
## We can drop these as follows:
flights %>%
  select(air_time) %>%
  filter(!duplicated(air_time)) %>%
```

```

arrange(air_time) %>%
slice(15:20)

## # A tibble: 6 x 1
##   air_time
##   <dbl>
## 1      34
## 2      35
## 3      36
## 4      37
## 5      38
## 6      39

## Another way:
flights %>%
  select(air_time) %>%
  distinct() %>%
  arrange(air_time) %>%
  slice(15:20)

## # A tibble: 6 x 1
##   air_time
##   <dbl>
## 1      34
## 2      35
## 3      36
## 4      37
## 5      38
## 6      39

## We can start composing the verbs to write more complex queries:
## the following lists the carriers with the longest departure delays
## on christmas day.
flights %>%
  select(year, month, day, dep_delay, carrier) %>%
  filter(month == 12, day == 25) %>%
  select(carrier, dep_delay) %>%
  arrange(desc(dep_delay)) %>%
  slice(1:10)

## # A tibble: 10 x 2
##   carrier dep_delay
##   <chr>      <dbl>
## 1      EV        321
## 2      EV        251
## 3      DL        234
## 4      UA        198
## 5      DL        193
## 6      EV        189
## 7      B6        170
## 8      B6        147
## 9      DL        146
## 10     UA        125

## While it wouldn't be very efficient to store full names in the
## original table, since we've now reduced the table down to a

```

```
## much smaller sub-table, it would be useful to be able to
## see the full names of carriers. To do this, we can "join" the
## two tables. "inner_join()" finds matching entries in both tables
## and combines them into one.
```

```
flights %>%
  select(year, month, day, dep_delay, carrier) %>%
  filter(month == 12, day == 25) %>%
  select(carrier, dep_delay) %>%
  arrange(desc(dep_delay)) %>%
  slice(1:10) %>%
  inner_join(airlines, by = "carrier") %>%
  select(name, dep_delay)
```

```
## # A tibble: 10 x 2
##           name dep_delay
##           <chr>     <dbl>
## 1 ExpressJet Airlines Inc.    321
## 2 ExpressJet Airlines Inc.    251
## 3   Delta Air Lines Inc.    234
## 4   United Air Lines Inc.    198
## 5   Delta Air Lines Inc.    193
## 6 ExpressJet Airlines Inc.    189
## 7   JetBlue Airways    170
## 8   JetBlue Airways    147
## 9   Delta Air Lines Inc.    146
## 10  United Air Lines Inc.    125
```

```
## Exercises 2.4:
##
##
## 1. Which carriers had the shortest departure delays?
```

```
flights %>%
  select(carrier, dep_delay) %>%
  arrange(dep_delay) %>%
  inner_join(airlines) %>%
  select(name, dep_delay)
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 336,776 x 2
##           name dep_delay
##           <chr>     <dbl>
## 1   JetBlue Airways    -43
## 2   Delta Air Lines Inc.   -33
## 3 ExpressJet Airlines Inc.   -32
## 4   Delta Air Lines Inc.   -30
## 5 Frontier Airlines Inc.   -27
## 6   Envoy Air    -26
## 7 ExpressJet Airlines Inc.   -25
## 8   Envoy Air    -25
## 9 Endeavor Air Inc.   -24
## 10  JetBlue Airways    -24
## # ... with 336,766 more rows
```




Figure 8: `Thedplyr::mutate()` function

```
## 2. Which flights had the longest flight (i.e. greatest distance)?
flights %>%
  select(carrier, tailnum, distance) %>%
  arrange(desc(distance)) %>%
  inner_join(airlines) %>%
  select(name, tailnum, distance)
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 336,776 x 3
```

```
##           name tailnum distance
##           <chr>   <chr>    <dbl>
## 1 Hawaiian Airlines Inc. N380HA      4983
## 2 Hawaiian Airlines Inc. N380HA      4983
## 3 Hawaiian Airlines Inc. N380HA      4983
## 4 Hawaiian Airlines Inc. N384HA      4983
## 5 Hawaiian Airlines Inc. N381HA      4983
## 6 Hawaiian Airlines Inc. N385HA      4983
## 7 Hawaiian Airlines Inc. N385HA      4983
## 8 Hawaiian Airlines Inc. N389HA      4983
## 9 Hawaiian Airlines Inc. N384HA      4983
## 10 Hawaiian Airlines Inc. N388HA      4983
## # ... with 336,766 more rows
```

4.1.4 `mutate()`

We can also derive new columns within a data frame by *mutating* existing columns with a call to `mutate()`. For example, it might be useful to have a column with information about how much time flights gained in the air, by computing the difference between the arrival and departure delays.

```
## Add the "gain" column described above.
```

```
flights %>%
  filter(!is.na(arr_delay), !is.na(dep_delay)) %>%
  mutate(gain = arr_delay - dep_delay,
         gain_per_hour = gain / (air_time / 60)) %>%
  arrange(desc(gain), desc(gain_per_hour))
```

```
## # A tibble: 327,346 x 21
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1 2013 11 1 658 700 -2 1329
## 2 2013 4 18 558 600 -2 1149
## 3 2013 8 8 1819 1519 180 5
## 4 2013 7 10 1916 1900 16 137
## 5 2013 6 27 1608 1525 43 2045
## 6 2013 7 22 1606 1615 -9 2056
## 7 2013 7 1 811 800 11 1344
## 8 2013 7 22 1626 1545 41 2051
## 9 2013 7 10 2011 1520 291 2357
## 10 2013 7 7 1659 1700 -1 2050
## # ... with 327,336 more rows, and 14 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, gain <dbl>, gain_per_hour <dbl>
```

It is also possible to discard all the old variables:

```
flights %>%
  filter(!is.na(arr_delay), !is.na(dep_delay)) %>%
  transmute(gain = arr_delay - dep_delay,
            gain_per_hour = gain / (air_time / 60)) %>%
  arrange(desc(gain), desc(gain_per_hour))
```

```
## # A tibble: 327,346 x 2
##   gain gain_per_hour
##   <dbl>         <dbl>
## 1 196      35.00000
## 2 181      46.41026
## 3 165      23.57143
## 4 161      24.96124
## 5 157      78.50000
## 6 154      66.00000
## 7 153      56.31902
## 8 150      93.75000
## 9 150      58.44156
## 10 148     138.75000
## # ... with 327,336 more rows
```

Sort in descending order:

```
flights %>%
  filter(!is.na(arr_delay), !is.na(dep_delay)) %>%
  mutate(gain = arr_delay - dep_delay,
         gain_per_hour = gain / (air_time / 60)) %>%
  arrange(desc(gain), desc(gain_per_hour)) %>%
  inner_join(airlines, by = "carrier") %>%
  select(name, gain, gain_per_hour)
```

```
## # A tibble: 327,346 x 3
##   name gain gain_per_hour
##   <chr> <dbl>         <dbl>
## 1 Virgin America 196      35.00000
## 2 American Airlines Inc. 181      46.41026
## 3 United Air Lines Inc. 165      23.57143
## 4 Delta Air Lines Inc. 161      24.96124
## 5 Envoy Air 157      78.50000
## 6 Delta Air Lines Inc. 154      66.00000
## 7 Delta Air Lines Inc. 153      56.31902
```

```
## 8          Envoy Air    150      93.75000
## 9 ExpressJet Airlines Inc. 150      58.44156
## 10         US Airways Inc. 148      138.75000
## # ... with 327,336 more rows
```

```
## All the "_time" columns in "flights" aren't in a very useful format.
## Let's convert these into "minutes since midnight" format.
```

```
flights %>%
  select(ends_with("_time"))
```

```
## # A tibble: 336,776 x 5
##   dep_time sched_dep_time arr_time sched_arr_time air_time
##   <int>      <int>      <int>      <int>      <dbl>
## 1      517          515        830          819        227
## 2      533          529        850          830        227
## 3      542          540        923          850        160
## 4      544          545       1004         1022        183
## 5      554          600        812          837        116
## 6      554          558        740          728        150
## 7      555          600        913          854        158
## 8      557          600        709          723         53
## 9      557          600        838          846        140
## 10     558          600        753          745        138
## # ... with 336,766 more rows
```

```
517 %/% 100      # Hour      (Integer division)
```

```
## [1] 5
```

```
517 %/% 100 * 60 # Hour in mins (Integer division)
```

```
## [1] 300
```

```
517 %% 100      # Minute      (Modulo remainder)
```

```
## [1] 17
```

```
## Let's write a function that does this for us.
```

```
conv_time = function(time) time %/% 100 * 60 + time %% 100
```

```
## 5:17 am means 5 hrs and 17 mins, or 317 mins since midnight.
```

```
conv_time(517)
```

```
## [1] 317
```

```
flights %>%
  mutate(dep_time_mins = conv_time(dep_time))
```

```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>      <int>      <dbl>   <int>
## 1  2013     1     1     517          515         2     830
## 2  2013     1     1     533          529         4     850
## 3  2013     1     1     542          540         2     923
## 4  2013     1     1     544          545        -1    1004
## 5  2013     1     1     554          600        -6     812
## 6  2013     1     1     554          558        -4     740
## 7  2013     1     1     555          600        -5     913
## 8  2013     1     1     557          600        -3     709
```

```
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 336,766 more rows, and 13 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, dep_time_mins <dbl>
```

```
## Exercise 2.5:
```

```
##
```

```
## 1. Create a mutated version of all columns that end with "_time" and
##   arrange them in descending order of the new "dep_time".
```

```
conv_time = function(time) time %/% 100 * 60 + time %% 100
```

```
flights %>%
  select(ends_with("_time")) %>%
  names()
```

```
## [1] "dep_time"      "sched_dep_time" "arr_time"      "sched_arr_time"
## [5] "air_time"
```

```
flights %>%
  mutate(dep_time2 = conv_time(dep_time),
         sched_dep_time2 = conv_time(sched_dep_time),
         arr_time2 = conv_time(arr_time),
         sched_arr_time2 = conv_time(sched_arr_time),
         air_time2 = conv_time(air_time)) %>%
  select(dep_time, dep_time2)
```

```
## # A tibble: 336,776 x 2
```

```
##   dep_time dep_time2
```

```
##   <int>      <dbl>
```

```
## 1     517      317
```

```
## 2     533      333
```

```
## 3     542      342
```

```
## 4     544      344
```

```
## 5     554      354
```

```
## 6     554      354
```

```
## 7     555      355
```

```
## 8     557      357
```

```
## 9     557      357
```

```
## 10    558      358
```

```
## # ... with 336,766 more rows
```

```
## 2. Modify the "conv_time" function above so that it returns the
##   time since midnight in HOURS. Then repeat Q1 above.
```

```
conv_time = function(time, hour = TRUE) {
  ifelse(hour,
         time %/% 100 + time %% 100 / 60,
         time %/% 100 * 60 + time %% 100)
}
```

```
conv_time(517)
```

```
## [1] 5.283333
```

```
conv_time(517, hour = FALSE)
```





Figure 9: The `dplyr::summarise()` function

```
## [1] 317
## 3. Does air_time equal arr_time - dep_time? Why or why not?
## (HINT: add a new column, and add another column computing
## the difference between the new column and air_time.)
flights %>%
  mutate(diff = mean(arr_time - dep_time, na.rm = TRUE),
         air_diff = mean(air_time - diff, na.rm = TRUE)) %>%
  filter(air_diff == 0)

## # A tibble: 0 x 21
## # ... with 21 variables: year <int>, month <int>, day <int>,
## #   dep_time <int>, sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   diff <dbl>, air_diff <dbl>

flights %>%
  transmute(diff = mean(arr_time - dep_time, na.rm = TRUE),
           air_diff = mean(air_time - diff, na.rm = TRUE)) %>%
  group_by(diff) %>%
  summarize(mean_air_diff = mean(air_diff))

## # A tibble: 1 x 2
##   diff mean_air_diff
##   <dbl>         <dbl>
## 1 153.1973      -2.510828
```

4.1.5 summarise()

The last of the five main `dplyr` verbs is `summarise`, which *collapses* a data frame down to a single row by applying a summary operation.

```
## Compute the mean arrival and departure delays, and drop missing values.
```

```
flights %>%
  summarize(mean_arr_delay = mean(arr_delay, na.rm = TRUE),
            mean_dep_delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 2
```

```
##   mean_arr_delay mean_dep_delay
##         <dbl>         <dbl>
## 1      6.895377      12.63907
```

summarise is particularly useful when applied to *groups* of observations by using it in conjunction with group_by().

```
## Mean arrival and departure delay times by carrier.
```

```
flights %>%
  filter(!is.na(arr_delay), !is.na(dep_delay)) %>%
  group_by(carrier) %>%
  summarize(mean_arr_delay = mean(arr_delay),
            mean_dep_delay = mean(dep_delay),
            count = n()) %>%
  inner_join(airlines) %>%
  select(name, mean_arr_delay, mean_dep_delay, count)
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 16 x 4
```

```
##           name mean_arr_delay mean_dep_delay count
##           <chr>         <dbl>         <dbl> <int>
## 1 Endeavor Air Inc.      7.3796692      16.439574 17294
## 2 American Airlines Inc.  0.3642909       8.569130 31947
## 3 Alaska Airlines Inc.  -9.9308886       5.830748  709
## 4 JetBlue Airways      9.4579733      12.967548 54049
## 5 Delta Air Lines Inc.   1.6443409       9.223950 47658
## 6 ExpressJet Airlines Inc. 15.7964311      19.838929 51108
## 7 Frontier Airlines Inc. 21.9207048      20.201175  681
## 8 AirTran Airways Corporation 20.1159055      18.605984  3175
## 9 Hawaiian Airlines Inc. -6.9152047       4.900585  342
## 10 Envoy Air      10.7747334      10.445381 25037
## 11 SkyWest Airlines Inc.  11.9310345      12.586207   29
## 12 United Air Lines Inc.  3.5580111      12.016908 57782
## 13 US Airways Inc.    2.1295951       3.744693 19831
## 14 Virgin America    1.7644644      12.756646  5116
## 15 Southwest Airlines Co. 9.6491199      17.661657 12044
## 16 Mesa Airlines Inc.  15.5569853      18.898897  544
```

```
## Exercise 2.6:
```

```
##
```

```
## 1. Which carrier had the lowest mean arrival delay?
```

```
flights_carr = flights %>%
  filter(!is.na(arr_delay), !is.na(dep_delay)) %>%
  group_by(carrier)
```

```
flights_carr %>%
  summarize(mean_arr_delay = mean(arr_delay)) %>%
  inner_join(airlines) %>%
  select(name, mean_arr_delay) %>%
```

```
arrange(mean_arr_delay)
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 16 x 2
```

```
##           name mean_arr_delay
##           <chr>         <dbl>
## 1   Alaska Airlines Inc.    -9.9308886
## 2   Hawaiian Airlines Inc.  -6.9152047
## 3   American Airlines Inc.   0.3642909
## 4   Delta Air Lines Inc.     1.6443409
## 5   Virgin America          1.7644644
## 6   US Airways Inc.         2.1295951
## 7   United Air Lines Inc.    3.5580111
## 8   Endeavor Air Inc.       7.3796692
## 9   JetBlue Airways         9.4579733
## 10  Southwest Airlines Co.   9.6491199
## 11   Envoy Air              10.7747334
## 12  SkyWest Airlines Inc.   11.9310345
## 13   Mesa Airlines Inc.     15.5569853
## 14  ExpressJet Airlines Inc. 15.7964311
## 15 AirTran Airways Corporation 20.1159055
## 16  Frontier Airlines Inc.  21.9207048
```

```
## 2. Which carrier had the lowest mean departure delay on xmas day?
```

```
flights_carr %>%
  filter(month == 12, day == 25) %>%
  summarize(mean_dep_delay = mean(dep_delay)) %>%
  inner_join(airlines) %>%
  select(name, mean_dep_delay) %>%
  arrange(mean_dep_delay)
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 14 x 2
```

```
##           name mean_dep_delay
##           <chr>         <dbl>
## 1   Alaska Airlines Inc.    -7.0000000
## 2   Hawaiian Airlines Inc.  -6.0000000
## 3   Frontier Airlines Inc.  -4.0000000
## 4 AirTran Airways Corporation -0.7142857
## 5   Envoy Air               1.0181818
## 6   US Airways Inc.         1.4166667
## 7   American Airlines Inc.   2.0384615
## 8   Virgin America          3.4615385
## 9   JetBlue Airways         3.8930818
## 10  Southwest Airlines Co.   6.7000000
## 11   Endeavor Air Inc.       7.0937500
## 12  United Air Lines Inc.   12.1570248
## 13   Delta Air Lines Inc.   14.4423077
## 14  ExpressJet Airlines Inc. 14.6933333
```

```
## 3. Which carrier had the highest maximum departure delay?
```

```
flights_carr %>%
  summarize(max_dep_delay = max(dep_delay, na.rm = TRUE)) %>%
  inner_join(airlines) %>%
```

```
select(name, max_dep_delay) %>%
  arrange(desc(max_dep_delay))
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 16 x 2
```

```
##           name max_dep_delay
##           <chr>         <dbl>
## 1 Hawaiian Airlines Inc.    1301
## 2 Envoy Air                 1137
## 3 American Airlines Inc.    1014
## 4 Delta Air Lines Inc.      960
## 5 Frontier Airlines Inc.    853
## 6 Endeavor Air Inc.        747
## 7 Virgin America           653
## 8 AirTran Airways Corporation 602
## 9 ExpressJet Airlines Inc.   548
## 10 JetBlue Airways          502
## 11 US Airways Inc.          500
## 12 United Air Lines Inc.     483
## 13 Southwest Airlines Co.    471
## 14 Mesa Airlines Inc.       387
## 15 Alaska Airlines Inc.     225
## 16 SkyWest Airlines Inc.    154
```

```
## 4. Which flights were the most delayed, based on arrival, on xmas day?
```

```
flights %>%
  filter(month == 12, day == 25) %>%
  inner_join(airlines) %>%
  select(name, arr_delay) %>%
  arrange(desc(arr_delay))
```

```
## Joining, by = "carrier"
```

```
## # A tibble: 719 x 2
```

```
##           name arr_delay
##           <chr>         <dbl>
## 1 ExpressJet Airlines Inc.    292
## 2 ExpressJet Airlines Inc.    250
## 3 Delta Air Lines Inc.        222
## 4 JetBlue Airways            194
## 5 Delta Air Lines Inc.        182
## 6 ExpressJet Airlines Inc.    180
## 7 United Air Lines Inc.       175
## 8 JetBlue Airways            139
## 9 Delta Air Lines Inc.        120
## 10 United Air Lines Inc.      117
## # ... with 709 more rows
```

```
## 5. Which season had the highest proportion of cancelled flights?
```

```
## Use the code below to get started.
```

```
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
```

```
##
##   date
## Function for adding seasonal information.
seas2013 = function(dates) {
  seasons = character(length(dates))
  seas = ymd(c("20130320", "20130621", "20130922", "20131221"))
  labs = c("Spring", "Summer", "Fall", "Winter")
  seasons[dates >= seas[4] | dates < seas[1]] = labs[4]
  for (i in 1:3) {
    seasons[dates >= seas[i] & dates < seas[i + 1]] = labs[i]
  }
  seasons
}

## Create a new version of the "flights" table with seasons.
flights2 = flights %>%
  filter(!is.na(year), !is.na(month), !is.na(day)) %>%
  mutate(date = make_date(year, month, day),
         seas = seas2013(date))

## The following gives the proportion of cancelled flights by month.
flights2 %>%
  filter(is.na(dep_time)) %>%
  group_by(month) %>%
  summarize(count = n()) %>%
  mutate(perc = round(count / sum(count) * 100, 2))
```

```
## # A tibble: 12 x 3
##   month count  perc
##   <int> <int> <dbl>
## 1     1    521  6.31
## 2     2   1261 15.28
## 3     3    861 10.43
## 4     4    668  8.09
## 5     5    563  6.82
## 6     6   1009 12.22
## 7     7    940 11.39
## 8     8    486  5.89
## 9     9    452  5.48
## 10    10    236  2.86
## 11    11    233  2.82
## 12    12   1025 12.42
```

```
## So back to the question:
##
## Which month had the highest proportion of cancelled flights?
flights2 %>%
  filter(is.na(dep_time)) %>%
  group_by(month) %>%
  summarize(count = n()) %>%
  mutate(perc = round(count / sum(count) * 100, 2)) %>%
  arrange(desc(perc))
```

```
## # A tibble: 12 x 3
##   month count  perc
```

```
##      <int> <int> <dbl>
## 1      2  1261 15.28
## 2     12  1025 12.42
## 3      6  1009 12.22
## 4      7   940 11.39
## 5      3   861 10.43
## 6      4   668  8.09
## 7      5   563  6.82
## 8      1   521  6.31
## 9      8   486  5.89
## 10     9   452  5.48
## 11    10   236  2.86
## 12    11   233  2.82
```

```
## Another one: Which season had the lowest proportion of cancelled flights?
```

```
flights2 %>%
  filter(is.na(dep_time)) %>%
  group_by(seas) %>%
  summarize(count = n()) %>%
  mutate(perc = round(count / sum(count) * 100, 2)) %>%
  arrange(perc)
```

```
## # A tibble: 4 x 3
##   seas count  perc
##   <chr> <int> <dbl>
## 1  Fall  1434 17.37
## 2 Spring 1854 22.46
## 3 Summer 2363 28.63
## 4 Winter 2604 31.54
```

4.2 Data Visualization with ggplot2

So far, we looked at how to program in R and learned how to manipulate data frames using the `dplyr` package. But all our work has been textual so far – we haven’t produced any pretty pictures! In this section, we will learn how to visualize data using the `ggplot2` package.

`ggplot2` allows us to build statistical graphics incrementally by providing a kind of *grammar* for data visualization, much like how `dplyr` provided a *grammar* for data manipulation. Not surprisingly, both packages were created by the same person – Chief Scientist of RStudio and Adjunct Professor of Statistics at the University of Auckland, Stanford University, and Rice University, Dr. Hadley Wickham. Whilst packages like `ggplot2` and `dplyr` do *not* add much in the way of *new* functionality to R, it does provide a uniform, user-friendly interface for interacting with data.

There are three key components in a plot produced using `ggplot2`:


- A set of **data** to plot,
- A set of **aesthetic** mappings, and
- A visualization layer.

An aesthetic is mapping between a variable and visual cues that can represent variable values, e.g. glyphs. A *glyph* is a basic graphical element that represents a single observation, e.g. mark, symbol. For example, in a scatterplot, the *positions* of a glyph on the plot act as visual cues that help us understand how two variables are related.


```
library(ggplot2)
```

```
## The "women" data set contains the heights and weights of 15 American
## women in ordinary indoor clothing and with shoes on. We can study
## the relationship between the variables "height" and "weight" by
## mapping the height to the x-axis and weight to the y-axis. The
## `aes()` function *maps* variables from the data to graphical
## attributes. We can assign this graphic to the variable "p".
p = ggplot(data = women, aes(x = height, y = weight))

## At this point, "p" is just an empty plot, as we haven't added a
## visualization layer. We can add features to "p" using the `+`
## operator. In this case, we add points.
p + geom_point()
```




```
p + geom_point(size = 3)
```



```
## We can also plot using pipe operators. Here, we convert heights and
## weights to sensible units and derive a new categorical variable
## called "weight_class".
women2 = women %>%
  mutate(height_cm = height * 2.54,
         weight_kg = weight * 0.453,
         weight_class = cut(weight, 3,
                           labels = c("Low", "Medium", "High")))


## Map the colour of each dot to "weight_class". Observe that we
## can omit specifying the x and y arguments in ggplot().
women2 %>%
  ggplot(aes(height, weight)) + # x-axis, y-axis
  geom_point(aes(color = weight_class))
```

```
## Map the size of each dot to "weight_class".
```


```
women2 %>%  
  ggplot(aes(height, weight)) +  
  geom_point(aes(size = weight_class))
```

```
## Warning: Using size for a discrete variable is not advised.
```




```
## The two examples above define an additional aesthetic (colour and size)
## to display the additional variable "weight_class". An alternative way
## is to use multiple side-by-side graphs, called "facets".
p = women2 %>%
  ggplot(aes(height, weight))

p + geom_point() +
  facet_wrap(~ weight_class)
```




```
## Exercise 3.1:
##
## 1. Derive a new variable called "height_class", with labels "Short",
##    "Average", and "Tall"; then produce a plot as above, but map the
##    color of each dot to the variable "height_class".
women2 = women %>%
  mutate(height_cm = height * 2.54,
         weight_kg = weight * 0.453,
         weight_class = cut(weight, 3,
                           labels = c("Low", "Medium", "High")),
         height_class = cut(height, 3,
                           labels = c("Short", "Average", "Tall")))

women2 %>%
  ggplot(aes(height, weight)) +
  geom_point(aes(col = height_class))
```




```
## 2. Use facets instead.  
women2 %>%  
  ggplot(aes(height, weight)) +  
  geom_point(aes(colour = height_class)) +  
  facet_wrap(~ height_class)
```




Other plots include:

```
## Line plot
women2 %>%
  ggplot(aes(height_cm, weight_kg)) +
  geom_line()
```



```
## Box plot
women2 %>%
  ggplot(aes(weight_class, height_cm)) +
  geom_boxplot(fill = "darkseagreen")
```



Now that we've covered the basics, let's visualize the `flights` data. Before we proceed any further though, we create an enriched version of the table to make it easier to visualize it over time.


```
## Read in enriched flights table.
flights2 = readRDS("flights2.Rds")
```

We are now ready to start visualizing the `flights` data. Let's start by plotting a histogram of the departure time. Recall that histograms divide up the x-axis into *bins* and count the number of observations that fall into each bin; it then displays the count with bars. We can adjust the `binwidth` to cover different ranges to reveal different characteristics of the data. It is also possible to specify the number of `bins`.

```
## Base plot for "dep_min".
p = flights2 %>% ggplot(aes(dep_min))


## Histogram of departure times in minutes.
p + geom_histogram(bins = 100)
```

```
## Warning: Removed 8255 rows containing non-finite values (stat_bin).
```



```
p + geom_histogram(binwidth = 100)
```


```
## Warning: Removed 8255 rows containing non-finite values (stat_bin).
```

```
## Lines instead of bars.
```


```
p + geom_freqpoly(binwidth = 100)
```

```
## Warning: Removed 8255 rows containing non-finite values (stat_bin).
```




```
## Distribution of xmas day flights.  
flights2 %>%  
  filter(month == 12, day == 25) %>%  
  ggplot(aes(dep_datetime)) +  
  geom_freqpoly(bins = 100)
```


```
## Warning: Removed 4 rows containing non-finite values (stat_bin).
```




```
## Number of flights by day of week.  
flights2 %>%  
  ggplot(aes(wday)) +  
  geom_bar(fill = "lightblue")
```




```
## Fewer flights on weekends.  
flights2 %>%  
  group_by(date) %>%  
  summarise(count = n()) %>%  
  mutate(wday = wday(date, lab = TRUE)) %>%  
  ggplot(aes(wday, count)) +  
  geom_boxplot()
```



```
## Exercise 3.2: Plot the distribution of flights (by weekday) in December.  
flights2 %>%  
  filter(month == 12) %>%  
  ggplot(aes(wday)) +  
  geom_bar(fill = "darkseagreen")
```




```
flights2 %>%  
  ggplot(aes(wday)) +  
  geom_bar(fill = "darkseagreen") +  
  facet_wrap(~ month)
```




4.2.1 Others

```
## Flights by hour
flights2 %>%
  ggplot(aes(hour)) +
  geom_bar() +
  scale_x_continuous(breaks = 0:24)
```




```
## Average departure delay among non-cancelled flights.  
## Observe that variation decreases with the number of flights.  
flights2 %>%  
  filter(!is.na(dep_delay)) %>%  
  group_by(tailnum) %>%  
  summarize(mean_delay = mean(dep_delay, na.rm = TRUE),  
            count = n()) %>%  
  ggplot(aes(count, mean_delay)) +  
  geom_point(alpha = 0.1) +  
  geom_smooth() # add model fit.  
  
## `geom_smooth()` using method = 'gam'
```


```
## Number of flights per day.
p2 = flights2 %>%
  group_by(date) %>%
  summarize(count = n()) %>%
  ggplot(aes(date, count)) +
  geom_line() +
  xlab("Date") +
  ylab("Number of Flights Scheduled")
```

p2




```
p2 + geom_smooth()
```

```
## `geom_smooth()` using method = 'loess'
```




```
p2 + geom_smooth(method = "lm")
```

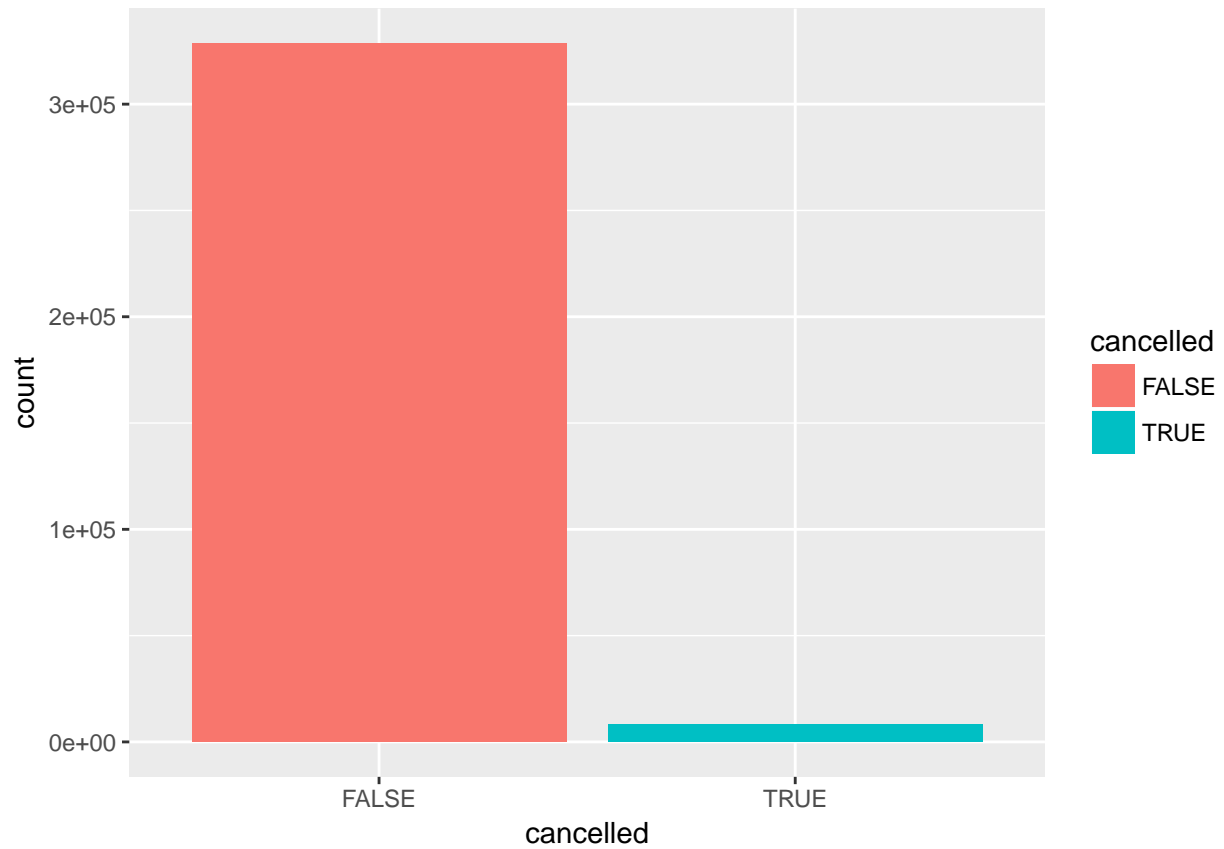


4.2.2 Cancelled flights

```
## Compare scheduled departure times for cancelled and non-cancelled.
flights2 %>%
  ggplot(aes(sched_dep_hr)) +
  geom_freqpoly(aes(colour = cancelled), binwidth = 0.5) +
  xlab("Hour") +
  ylab("Number of Flights Scheduled")
```




```
flights2 %>%  
  ggplot(aes(cancelled)) +  
  geom_bar(aes(fill = cancelled))
```




4.2.3 Delayed flights

```
## Departure delay
flights2 %>%
  filter(!is.na(dep_delayed)) %>%
  ggplot(aes(dep_delayed)) +
  geom_bar(aes(fill = dep_delayed))
```



```
## Arrival delay
flights2 %>%
  filter(!is.na(arr_delayed)) %>%
  ggplot(aes(arr_delayed)) +
  geom_bar(aes(fill = arr_delayed))
```



```
## Departure delay by airport of origin
flights2 %>%
  group_by(origin) %>%
  summarise(mean_dep_delay = mean(dep_delay, na.rm = TRUE)) %>%
  ggplot(aes(origin, mean_dep_delay)) +
  geom_col() # geom_bar() uses stat = count() by default.
```

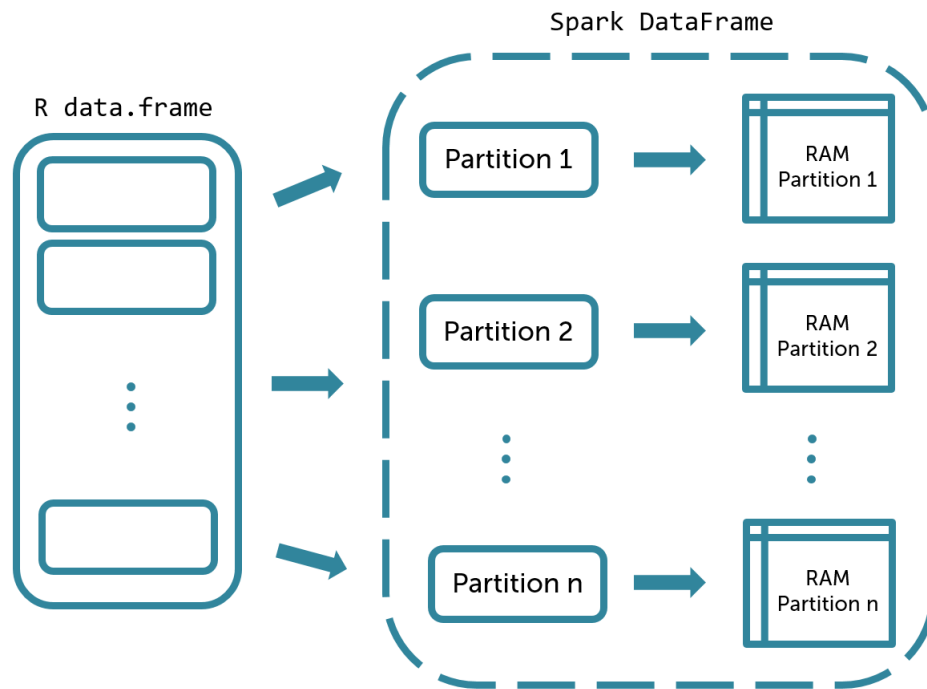




Figure 10: Distributing a local R `data.frame` to a Spark `DataFrame`






Figure 11: Row-based vs. column-based Storage

4.3 “Distributed Processing” with sparklyr

4.3.1 Setup:

```
## Load required packages.
library(dplyr)
library(sparklyr)

##
## Attaching package: 'sparklyr'

## The following object is masked from 'package:dplyr':
##
##   top_n

library(nycflights13)
library(ggplot2)
library(lubridate)

## Set config parameters for training laptops.
config = spark_config()
config$spark.executor.cores = 4      # parallel::detectCores()
config$spark.executor.memory = "4G"  # approx. .5 * memory.size()
config$spark.driver.memory = "4G"

## Connect to Spark cluster
sc = spark_connect("local", config = config)
# spark_disconnect(sc)
```

4.3.2 Efficient storage

We store the table as a Parquet file, which is an efficient column-based format for storing data. Each column is stored with other columns of the same type, and data is spread across different blocks on disk.

Why is this a good thing? suppose we want to **filter** the rows of `flights` based on the condition `dep_delay > 120`, i.e. departure delays of more than 2 hours. If the data is stored in a record-based format like `.csv`, the query needs to:

- Scan and read in the row,
- Parse the row into different columns, e.g. `year`, `month`, `day`, `dep_delay`, etc.,
- Obtain the relevant column (`dep_delay`), and
- Filter the row based on the condition `dep_delay > 120`.

for *every* row. If we use a column-based format like Parquet instead, the query only reads in the column of interest (`dep_delay`) - it ignores the other columns as it knows where that column resides and what type it is, all in advance. This results in a lot less work.

```
## Read in modified flights data.
flights2_rdf = readRDS("flights2_rdf.Rds")

## Create a Spark DataFrame out of the local R data.frame.
# flights_sdf = copy_to(sc, flights2_rdf, "flights", overwrite = TRUE)

## Save as Parquet file. Parquet is a columnar file format, cf.
## csv which is record or row-based. When querying the latter,
## the query needs to scan every row of the data set: read it in,
## parse it into fields, etc. With columnar formats, the query
## is only concerned with the columns required for the query,
## and doesn't bother with the others - so we can avoid doing
## a lot of work.
# spark_write_parquet(flights_sdf, "flights_sdf.parquet")

## Read flights table and register as Spark DataFrame.
flights_sdf = spark_read_parquet(sc, "flights_sdf", "flights_sdf.parquet") %>%
  sdf_register(name = "flights")

## Have an initial look at the table.
head(flights_sdf)

## # Source:   lazy query [?? x 43]
## # Database: spark_connection
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>   <int>
## 1  2013    12     5     1800           1706          54     2104
## 2  2013    12     5     1802           1535         147     1933
## 3  2013    12     5     1804           1645          79     1922
## 4  2013    12     5     1806           1745          21     2013
## 5  2013    12     5     1807           1810          -3     2134
## 6  2013    12     5     1807           1759           8     2120
## # ... with 36 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <int>, minute <int>,
## #   cancelled <lgl>, dep_delayed <lgl>, arr_delayed <lgl>, delayed <lgl>,
## #   seas <chr>, wday <chr>, week <dbl>, month_lab <chr>, id <int>,
## #   dep_min <dbl>, dep_hr <dbl>, sched_dep_min <dbl>, sched_dep_hr <dbl>,
## #   arr_min <dbl>, arr_hr <dbl>, sched_arr_min <dbl>, sched_arr_hr <dbl>,
## #   air_time_min <dbl>, air_time_hr <dbl>, date <chr>, dep_datetime <chr>,
## #   arr_datetime <chr>, sched_dep_datetime <chr>,
## #   sched_arr_datetime <chr>, time_hour <chr>

## How many partitions?
sdf_num_partitions(flights_sdf)

## [1] 8
```

4.3.3 Caching

In general, storing data in RAM is faster than disk (e.g. Hard Drive) because RAM sits very close to the CPU (the *brain* of the computer) so that data has to travel less. Moreover, physical hard disks have mechanical parts that move around and find a place to store data (think of a needle on an old record player finding a song of your choice), which takes time. RAM stands for Random Access Memory, because computers are said to be able to access any region of memory equally quickly. Disk space is about 100 times cheaper per byte than RAM, but between 5-100,000 times slower depending on the type of workload.

Spark allows us to pull data sets into a cluster-wide in-memory cache. Cache is memory that is much closer to the CPU than RAM. This means transfers to and from cache takes much less time than RAM (roughly 10 times faster).

```
## Register as Spark DataFrame and give it a table name for the Spark SQL
## context.
# microbenchmark::microbenchmark(
#   flights_sdf %>%
#     sdf_sample(frac = 1000, replace = TRUE) %>%
#     filter(!is.na(dep_min), !is.na(arr_min)) %>%
#     group_by(wday) %>%
#     summarise(mean_dep_delay = mean(dep_min),
#               mean_arr_delay = mean(arr_min))
# )
#      min      lq     mean  median      uq     max  neval
# 46.21666 52.87225 68.25759 57.39077 75.1553 155.1505    100

## Cache table.
tbl_cache(sc, "flights")

## Exercise 3.3: Why did I do this (HINT: Actions and transformations).
head(flights_sdf)
```

```
## # Source:   lazy query [?? x 43]
## # Database: spark_connection
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>   <int>
## 1  2013    12     5    1800           1706          54    2104
## 2  2013    12     5    1802           1535         147    1933
## 3  2013    12     5    1804           1645          79    1922
## 4  2013    12     5    1806           1745          21    2013
## 5  2013    12     5    1807           1810          -3    2134
## 6  2013    12     5    1807           1759           8    2120
## # ... with 36 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <int>, minute <int>,
## #   cancelled <lgl>, dep_delayed <lgl>, arr_delayed <lgl>, delayed <lgl>,
## #   seas <chr>, wday <chr>, week <dbl>, month_lab <chr>, id <int>,
## #   dep_min <dbl>, dep_hr <dbl>, sched_dep_min <dbl>, sched_dep_hr <dbl>,
## #   arr_min <dbl>, arr_hr <dbl>, sched_arr_min <dbl>, sched_arr_hr <dbl>,
## #   air_time_min <dbl>, air_time_hr <dbl>, date <chr>, dep_datetime <chr>,
## #   arr_datetime <chr>, sched_dep_datetime <chr>,
## #   sched_arr_datetime <chr>, time_hour <chr>

# microbenchmark::microbenchmark(
#   flights_sdf %>%
#     filter(!is.na(dep_min), !is.na(arr_min)) %>%
```


```

#   group_by(wday) %>%
#   summarise(mean_dep_delay = mean(dep_min),
#             mean_arr_delay = mean(arr_min))
# )
#   min      lq    mean  median      uq    max  neval
# 4.380441 4.86123 6.506899 5.272292 5.665379 37.34042    100


```

4.3.4 Plots

```
## Run some example code. Observe that only difference is
## `collect()`ing the distributed Spark DataFrame into a
## local R data frame prior to plotting.
flights_sdf %>%
  select(year, month, day, dep_delay, arr_delay) %>%
  filter(!is.na(arr_delay)) %>%
  mutate(gain = dep_delay - arr_delay) %>%
  filter(gain > 60) %>%
  collect() %>% # Collect distributed structure into memory.
  mutate(date = make_date(year, month, day)) %>%
  ggplot(aes(date, gain)) +
  geom_point(alpha = 0.5, col = "darkseagreen", cex = 2)
```



```
flights_sdf %>%
  select(sched_dep_hr, cancelled) %>%
  collect() %>%
  ggplot(aes(sched_dep_hr)) +
  geom_freqpoly(aes(colour = cancelled), binwidth = 0.5) +
  xlab("Hour") +
  ylab("Number of Flights Scheduled")
```



```
## Exercise 3.4: Plot the distribution of flights (by week day) of ALL months
##               using sparklyr and ggplot2.
flights_sdf %>%
  select(wday, month) %>%
  collect() %>%
  ggplot(aes(wday)) +
  geom_bar(fill = "darkseagreen") +
  facet_wrap(~ month)
```

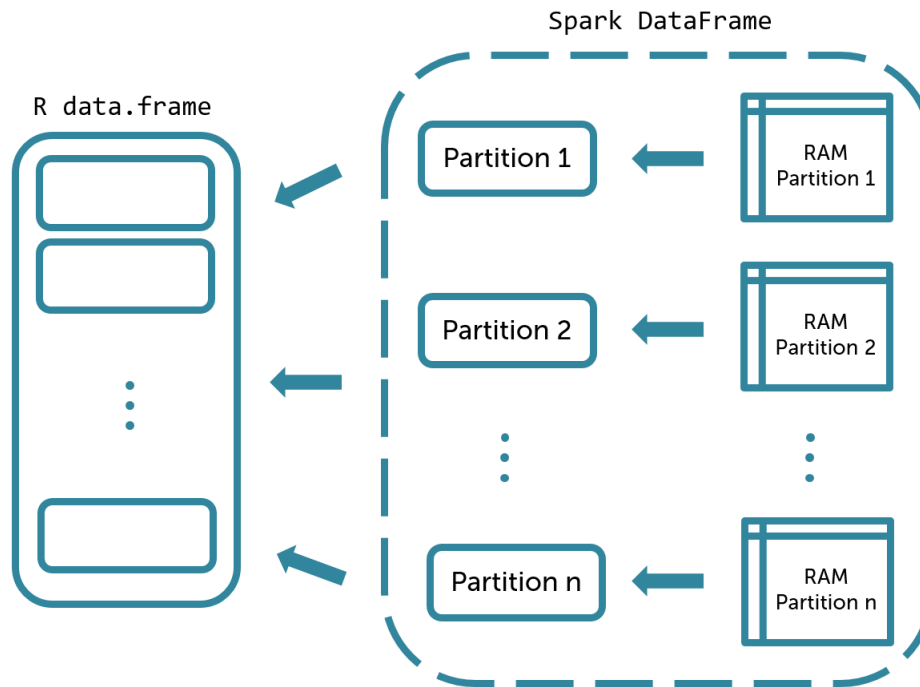



Figure 12: Collecting a Spark DataFrame to a local R data.frame



4.3.5 SQL

It is also possible to write queries in SQL against Spark tables.

```
library(DBI)
dbGetQuery(sc, "
SELECT dep_time FROM flights_sdf
WHERE dep_time > 100 AND NOT(dep_time) IS NULL
LIMIT 5
")

##    dep_time
## 1      1800
## 2      1802
## 3      1804
## 4      1806
## 5      1807

## It is in fact possible to translate "sparklyr" code directly to SQL.
q = flights_sdf %>%
  select(dep_time) %>%
  filter(dep_time > 100, !is.na(dep_time)) %>%
  head(5)

show_query(q)

## <SQL>
## SELECT *
## FROM (SELECT `dep_time` AS `dep_time`
## FROM `flights`) `nctmfrruvk`
## WHERE ((`dep_time` > 100.0) AND (NOT(((`dep_time`) IS NULL))))
## LIMIT 5

## Exercise 3.5: Plot the distribution of flights (by week day) of ALL months
##               for non-cancelled flights with a departure delay of more
##               than 1 hour using ggplot2, where:
##
## 1. Pure dplyr
q1 = flights2 %>%
  filter(!is.na(dep_delay), dep_delay > 60) %>%
  select(wday, month)

## 2. sparklyr (R code)
q2 = flights_sdf %>%
  filter(!is.na(dep_delay), dep_delay > 60) %>%
  select(wday, month)

## 3. SQL interface to spark using DBI.
show_query(q2) # Hint

## <SQL>
## SELECT `wday` AS `wday`, `month` AS `month`
## FROM `flights`
## WHERE ((NOT(((`dep_delay`) IS NULL))) AND (`dep_delay` > 60.0))

sql_q = paste("SELECT wday, month",
               "FROM flights_sdf",
```


```

      "WHERE NOT(dep_delay) IS NULL AND dep_delay > 60")
q3 = DBI::dbGetQuery(sc, sql_q)

## Plot
plot_q = function(x) {
  x %>%
    ggplot(aes(wday)) +
    geom_bar(fill = "darkseagreen") +
    facet_wrap(~ month)
}

plot_q(q1)


```




```

plot_q(q2 %>% collect())

```




```
plot_q(q3)
```



4.3.6 Models

```
## Create test and training sets.
flights_part = flights_sdf %>%
  filter(dep_delay > 240, arr_delay > 240,
         !is.na(dep_delay), !is.na(arr_delay)) %>%
  select(dep_delay, arr_delay) %>%
  sdf_partition(training = 0.7, test = 0.3, seed = 1234)

## Plot
flights_part$training %>%
  collect() %>%
  ggplot(aes(dep_delay, arr_delay)) +
  geom_line() +
  geom_smooth(method = "lm")
```




```
## Fit a linear model.
flights_lm = flights_part$training %>%
  ml_linear_regression(dep_delay ~ arr_delay)

## * No rows dropped by 'na.omit' call
## Look at model summary.
summary(flights_lm)

## Call: ml_linear_regression(., dep_delay ~ arr_delay)
##
## Deviance Residuals::
##      Min       1Q   Median       3Q      Max
## -148.658  -10.109    4.237   16.233   68.081
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept)  8.7152834   2.7152678   3.2097  0.001375 **
## arr_delay    0.9771936   0.0079103  123.5336 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-Squared: 0.944
## Root Mean Squared Error: 26.69
flights_lm$mean.absolute.error

## [1] 19.20164
```

```
## See predicted vs actual.
sdf_predict(flights_lm, flights_part$test) %>%
  collect() %>%
  ggplot(aes(dep_delay, prediction)) +
  geom_abline(lty = "dashed", col = "red") +
  geom_point(alpha = 0.1) +
  coord_fixed(ratio = 1) +
  labs(title = "Predicted vs Actual Extreme Departure Delays")
```



4.4 Review

The aim of today's workshop was to get hands-on experience in dealing with *reasonably large* data sets using R. You should now be able to:

- Read and write simple programs in the R language,
- Manipulate tabular data sets in R using the `dplyr` package,
- Visualize tabular data sets in R using the `ggplot2` package,
- Understand R's limitations when it comes to large data sets, and
- Manipulate *reasonably large* tabular data sets in R using the `sparklyr` package.

4.5 Appendix

4.5.1 Project Ideas

- [sparklyr Web Applications](#) to combine `nycflights13` with `leaflet`.
- [NYC Flights 2014 Data](#) for comparing with `nycflights13`.
- [R Data Sets](#)

4.5.2 Assignment Operator

Some of you may be more familiar with `<-` for variable assignment. In this course we use the C-style `=` operator instead for greater compatibility with mainstream programming languages (it also saves us from having to type an extra keystroke every time we assign a variable).

The main difference between the two operators is that `=` is only allowed in two places: at the top level, or within braces/extra pair of parens. For example, to assign a variable during a function call, one would have to wrap the expression within an extra pair of parens:

```
# mean(y = 1:10)      # illegal
# mean((y = 1:10))    # legal
# mean(y <- 1:10)     # legal
```

Another example might be in control structures:

```
# if (x = 0) 1 else x  # illegal
# if ((x = 0)) 1 else x # legal
# if (x <- 0) 1 else x  # legal
```

However, variable assignment during a function call or within control structures can cause confusion, especially when used named arguments. At the end of the day, the choice comes down to preference, but the course uses `=` to avoid using constructs that are unique to domain-specific programming languages like R. If you're interested, you can read more about assignment operators on the [R Developer website](#).

4.6 R References

- [Book: Modern Data Science with R](#)
- [Book: Parallel Computing for Data Science](#)
- [Book: R for Data Science](#)
- [Tutorial: Hexagonal Binning](#)
- [Tutorial: Introduction to `dplyr`](#)
- [Tutorial: Introduction to `ggplot2`](#)
- [Tutorial: `rcpp`](#)
- [Tutorial: Spatial Visualizations in R](#)
- [Website: Leaflet for R](#)
- [Website: `sparklyr`](#)
- [Website: `SparkR`](#)